

# **Class08\_Mini\_Project**

Qihao Liu

## **Table of contents**

<b>Exploratory data analysis</b>	<b>2</b>
Data import . . . . .	2
Exploratory data analysis . . . . .	4
<b>Principal Component Analysis</b>	<b>5</b>
Performing PCA . . . . .	5
Interpreting PCA results . . . . .	8
Variance Explained . . . . .	12
Communicating PCA results . . . . .	15
<b>Hierarchical clustering</b>	<b>16</b>
Results of hierarchical clustering . . . . .	16
Selecting number of clusters . . . . .	18
Using different methods . . . . .	21
<b>Optional: K-means Clustering</b>	<b>21</b>
k-means clustering and comparing results . . . . .	21
<b>Combining methods</b>	<b>22</b>
Clustering on PCA results . . . . .	22
<b>Sensitivity &amp; Specificity</b>	<b>24</b>
<b>Prediction</b>	<b>27</b>

## Exploratory data analysis

### Data import

Data was downloaded from class website as a .csv file. The csv file is imported into R as a data frame.

Note that the first column here wisc.df\$diagnosis is a pathologist provided expert diagnosis. We will not be using this for our unsupervised analysis as it is essentially the “answer” to the question which cell samples are malignant or benign.

```
# Save your input data file into your Project directory
fna.data <- read.csv("WisconsinCancer.csv")

# Complete the following code to input the data and store as wisc.df
wisc.df <- read.csv("WisconsinCancer.csv", row.names=1)

# We can use -1 here to remove the first column
wisc.data <- wisc.df[,-1]
head(wisc.data)
```

	radius_mean	texture_mean	perimeter_mean	area_mean	smoothness_mean
842302	17.99	10.38	122.80	1001.0	0.11840
842517	20.57	17.77	132.90	1326.0	0.08474
84300903	19.69	21.25	130.00	1203.0	0.10960
84348301	11.42	20.38	77.58	386.1	0.14250
84358402	20.29	14.34	135.10	1297.0	0.10030
843786	12.45	15.70	82.57	477.1	0.12780
	compactness_mean	concavity_mean	concave.points_mean	symmetry_mean	
842302	0.27760	0.3001	0.14710	0.2419	
842517	0.07864	0.0869	0.07017	0.1812	
84300903	0.15990	0.1974	0.12790	0.2069	
84348301	0.28390	0.2414	0.10520	0.2597	
84358402	0.13280	0.1980	0.10430	0.1809	
843786	0.17000	0.1578	0.08089	0.2087	
	fractal_dimension_mean	radius_se	texture_se	perimeter_se	area_se
842302	0.07871	1.0950	0.9053	8.589	153.40
842517	0.05667	0.5435	0.7339	3.398	74.08
84300903	0.05999	0.7456	0.7869	4.585	94.03
84348301	0.09744	0.4956	1.1560	3.445	27.23
84358402	0.05883	0.7572	0.7813	5.438	94.44
843786	0.07613	0.3345	0.8902	2.217	27.19

	smoothness_se	compactness_se	concavity_se	concave.points_se	
842302	0.006399	0.04904	0.05373	0.01587	
842517	0.005225	0.01308	0.01860	0.01340	
84300903	0.006150	0.04006	0.03832	0.02058	
84348301	0.009110	0.07458	0.05661	0.01867	
84358402	0.011490	0.02461	0.05688	0.01885	
843786	0.007510	0.03345	0.03672	0.01137	
	symmetry_se	fractal_dimension_se	radius_worst	texture_worst	
842302	0.03003	0.006193	25.38	17.33	
842517	0.01389	0.003532	24.99	23.41	
84300903	0.02250	0.004571	23.57	25.53	
84348301	0.05963	0.009208	14.91	26.50	
84358402	0.01756	0.005115	22.54	16.67	
843786	0.02165	0.005082	15.47	23.75	
	perimeter_worst	area_worst	smoothness_worst	compactness_worst	
842302	184.60	2019.0	0.1622	0.6656	
842517	158.80	1956.0	0.1238	0.1866	
84300903	152.50	1709.0	0.1444	0.4245	
84348301	98.87	567.7	0.2098	0.8663	
84358402	152.20	1575.0	0.1374	0.2050	
843786	103.40	741.6	0.1791	0.5249	
	concavity_worst	concave.points_worst	symmetry_worst		
842302	0.7119	0.2654	0.4601		
842517	0.2416	0.1860	0.2750		
84300903	0.4504	0.2430	0.3613		
84348301	0.6869	0.2575	0.6638		
84358402	0.4000	0.1625	0.2364		
843786	0.5355	0.1741	0.3985		
	fractal_dimension_worst				
842302	0.11890				
842517	0.08902				
84300903	0.08758				
84348301	0.17300				
84358402	0.07678				
843786	0.12440				

Setup a separate new vector called diagnosis that contains the data from the diagnosis column of the original dataset. We will store this as a factor (useful for plotting) and use this later to check our results. (we want to test clustering to see if it partitions the data correctly, having this would be cheating)

```
diagnosis <- factor(wisc.df$diagnosis)
head(diagnosis)
```

```
[1] M M M M M M
Levels: B M
```

## Exploratory data analysis

Q1. How many observations are in this dataset?

There are 569 observations in this dataset

Q2. How many of the observations have a malignant diagnosis?

There are 212 malignant diagnosis

```
table(wisc.df$diagnosis)
```

B	M
357	212

Q3. How many variables/features in the data are suffixed with \_mean?

There are 10 variables/features in the data are suffixed with “\_mean”

```
colnames(wisc.df)
```

```
[1] "diagnosis"                  "radius_mean"
[3] "texture_mean"               "perimeter_mean"
[5] "area_mean"                  "smoothness_mean"
[7] "compactness_mean"            "concavity_mean"
[9] "concave.points_mean"        "symmetry_mean"
[11] "fractal_dimension_mean"     "radius_se"
[13] "texture_se"                 "perimeter_se"
[15] "area_se"                    "smoothness_se"
[17] "compactness_se"              "concavity_se"
[19] "concave.points_se"           "symmetry_se"
[21] "fractal_dimension_se"       "radius_worst"
[23] "texture_worst"              "perimeter_worst"
[25] "area_worst"                 "smoothness_worst"
```

```
[27] "compactness_worst"           "concavity_worst"  
[29] "concave.points_worst"       "symmetry_worst"  
[31] "fractal_dimension_worst"
```

```
grep("_mean", colnames(wisc.df))
```

```
[1] 2 3 4 5 6 7 8 9 10 11
```

```
length(grep("_mean", colnames(wisc.df)))
```

```
[1] 10
```

## Principal Component Analysis

### Performing PCA

Note: use the `prcomp()` function to do PCA

- If we look at the argument in `prcomp()`, we can see that `scale.` = `FALSE` by default. However, if a column has a high value of variance, and `scale.` = `FALSE`, PCA will be dominated by columns with large variance (*not a feature, just the nature of numbers for certain parameter, i.e. speed of walking for different people is give or take ~4-5m/s, but the number of hairs is ~90k-150k. Number of hair would have a much larger variance, which will dominate PCA, but it doesn't necessarily makes it a more meaningful way to cluster different groups of people*). That's why **in general we should use `scale.` = `TRUE`**

Check the mean and standard deviation of the features (i.e. columns) of the `wisc.data` to determine if the data should be scaled. Use the `colMeans()` and `apply()` functions like you've done before.

```
# Check column means and standard deviations  
colMeans(wisc.data)
```

radius_mean	texture_mean	perimeter_mean
1.412729e+01	1.928965e+01	9.196903e+01
area_mean	smoothness_mean	compactness_mean
6.548891e+02	9.636028e-02	1.043410e-01
concavity_mean	concave.points_mean	symmetry_mean
8.879932e-02	4.891915e-02	1.811619e-01

fractal_dimension_mean		radius_se	texture_se
6.279761e-02		4.051721e-01	1.216853e+00
perimeter_se		area_se	smoothness_se
2.866059e+00		4.033708e+01	7.040979e-03
compactness_se		concavity_se	concave.points_se
2.547814e-02		3.189372e-02	1.179614e-02
symmetry_se	fractal_dimension_se		radius_worst
2.054230e-02		3.794904e-03	1.626919e+01
texture_worst	perimeter_worst		area_worst
2.567722e+01		1.072612e+02	8.805831e+02
smoothness_worst	compactness_worst	concavity_worst	
1.323686e-01		2.542650e-01	2.721885e-01
concave.points_worst	symmetry_worst	fractal_dimension_worst	
1.146062e-01		2.900756e-01	8.394582e-02

```
apply(wisc.data, 2, sd)
```

radius_mean	texture_mean	perimeter_mean
3.524049e+00	4.301036e+00	2.429898e+01
area_mean	smoothness_mean	compactness_mean
3.519141e+02	1.406413e-02	5.281276e-02
concavity_mean	concave.points_mean	symmetry_mean
7.971981e-02	3.880284e-02	2.741428e-02
fractal_dimension_mean	radius_se	texture_se
7.060363e-03	2.773127e-01	5.516484e-01
perimeter_se	area_se	smoothness_se
2.021855e+00	4.549101e+01	3.002518e-03
compactness_se	concavity_se	concave.points_se
1.790818e-02	3.018606e-02	6.170285e-03
symmetry_se	fractal_dimension_se	radius_worst
8.266372e-03	2.646071e-03	4.833242e+00
texture_worst	perimeter_worst	area_worst
6.146258e+00	3.360254e+01	5.693570e+02
smoothness_worst	compactness_worst	concavity_worst
2.283243e-02	1.573365e-01	2.086243e-01
concave.points_worst	symmetry_worst	fractal_dimension_worst
6.573234e-02	6.186747e-02	1.806127e-02

We can see that there's large difference between the standard deviations of different variables, so we should use scaling

Next, we do our PCA

```
# Perform PCA on wisc.data by completing the following code
wisc.pr <- prcomp(wisc.data,scale. = TRUE)

# Look at summary of results
summary(wisc.pr)
```

Importance of components:

	PC1	PC2	PC3	PC4	PC5	PC6	PC7
Standard deviation	3.6444	2.3857	1.67867	1.40735	1.28403	1.09880	0.82172
Proportion of Variance	0.4427	0.1897	0.09393	0.06602	0.05496	0.04025	0.02251
Cumulative Proportion	0.4427	0.6324	0.72636	0.79239	0.84734	0.88759	0.91010
	PC8	PC9	PC10	PC11	PC12	PC13	PC14
Standard deviation	0.69037	0.6457	0.59219	0.5421	0.51104	0.49128	0.39624
Proportion of Variance	0.01589	0.0139	0.01169	0.0098	0.00871	0.00805	0.00523
Cumulative Proportion	0.92598	0.9399	0.95157	0.9614	0.97007	0.97812	0.98335
	PC15	PC16	PC17	PC18	PC19	PC20	PC21
Standard deviation	0.30681	0.28260	0.24372	0.22939	0.22244	0.17652	0.1731
Proportion of Variance	0.00314	0.00266	0.00198	0.00175	0.00165	0.00104	0.0010
Cumulative Proportion	0.98649	0.98915	0.99113	0.99288	0.99453	0.99557	0.9966
	PC22	PC23	PC24	PC25	PC26	PC27	PC28
Standard deviation	0.16565	0.15602	0.1344	0.12442	0.09043	0.08307	0.03987
Proportion of Variance	0.00091	0.00081	0.0006	0.00052	0.00027	0.00023	0.00005
Cumulative Proportion	0.99749	0.99830	0.9989	0.99942	0.99969	0.99992	0.99997
	PC29	PC30					
Standard deviation	0.02736	0.01153					
Proportion of Variance	0.00002	0.00000					
Cumulative Proportion	1.00000	1.00000					

Q4. From your results, what proportion of the original variance is captured by the first principal components (PC1)?

44.27% of the original variance is captured by PC1

Q5. How many principal components (PCs) are required to describe at least 70% of the original variance in the data?

3 PCs are required, because the Cumulative Proportion for PC3 is 72.636%

Q6. How many principal components (PCs) are required to describe at least 90% of the original variance in the data?

7 PCs are required

## Interpreting PCA results

Now you will use some visualizations to better understand your PCA model.

First, we use gg plot to make scatter plot

```
library(ggplot2)
head(wisc.pr$x)
```

	PC1	PC2	PC3	PC4	PC5	PC6
842302	-9.184755	-1.946870	-1.1221788	3.6305364	1.1940595	1.41018364
842517	-2.385703	3.764859	-0.5288274	1.1172808	-0.6212284	0.02863116
84300903	-5.728855	1.074229	-0.5512625	0.9112808	0.1769302	0.54097615
84348301	-7.116691	-10.266556	-3.2299475	0.1524129	2.9582754	3.05073750
84358402	-3.931842	1.946359	1.3885450	2.9380542	-0.5462667	-1.22541641
843786	-2.378155	-3.946456	-2.9322967	0.9402096	1.0551135	-0.45064213
	PC7	PC8	PC9	PC10	PC11	PC12
842302	2.15747152	0.39805698	-0.15698023	-0.8766305	-0.2627243	-0.8582593
842517	0.01334635	-0.24077660	-0.71127897	1.1060218	-0.8124048	0.1577838
84300903	-0.66757908	-0.09728813	0.02404449	0.4538760	0.6050715	0.1242777
84348301	1.42865363	-1.05863376	-1.40420412	-1.1159933	1.1505012	1.0104267
84358402	-0.93538950	-0.63581661	-0.26357355	0.3773724	-0.6507870	-0.1104183
843786	0.49001396	0.16529843	-0.13335576	-0.5299649	-0.1096698	0.0813699
	PC13	PC14	PC15	PC16	PC17	
842302	0.10329677	-0.690196797	0.601264078	0.74446075	-0.26523740	
842517	-0.94269981	-0.652900844	-0.008966977	-0.64823831	-0.01719707	
84300903	-0.41026561	0.016665095	-0.482994760	0.32482472	0.19075064	
84348301	-0.93245070	-0.486988399	0.168699395	0.05132509	0.48220960	
84358402	0.38760691	-0.538706543	-0.310046684	-0.15247165	0.13302526	
843786	-0.02625135	0.003133944	-0.178447576	-0.01270566	0.19671335	
	PC18	PC19	PC20	PC21	PC22	
842302	-0.54907956	0.1336499	0.34526111	0.096430045	-0.06878939	
842517	0.31801756	-0.2473470	-0.11403274	-0.077259494	0.09449530	
84300903	-0.08789759	-0.3922812	-0.20435242	0.310793246	0.06025601	
84348301	-0.03584323	-0.0267241	-0.46432511	0.433811661	0.20308706	
84358402	-0.01869779	0.4610302	0.06543782	-0.116442469	0.01763433	
843786	-0.29727706	-0.1297265	-0.07117453	-0.002400178	0.10108043	
	PC23	PC24	PC25	PC26	PC27	
842302	0.08444429	0.175102213	0.150887294	-0.201326305	-0.25236294	
842517	-0.21752666	-0.011280193	0.170360355	-0.041092627	0.18111081	
84300903	-0.07422581	-0.102671419	-0.171007656	0.004731249	0.04952586	
84348301	-0.12399554	-0.153294780	-0.077427574	-0.274982822	0.18330078	

```

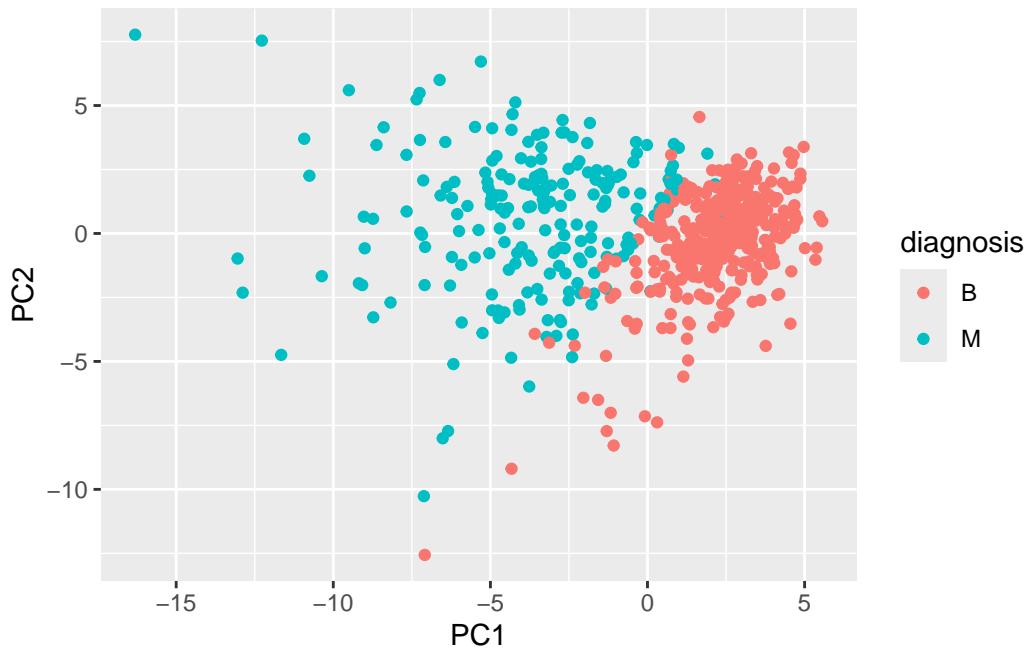
84358402  0.13933105  0.005327110 -0.003059371  0.039219780  0.03213957
843786    0.03344819 -0.002837749 -0.122282765 -0.030272333 -0.08438081
PC28      PC29      PC30
842302   -0.0338846387  0.045607590  0.0471277407
842517    0.0325955021 -0.005682424  0.0018662342
84300903  0.0469844833  0.003143131 -0.0007498749
84348301  0.0424469831 -0.069233868  0.0199198881
84358402 -0.0347556386  0.005033481 -0.0211951203
843786    0.0007296587 -0.019703996 -0.0034564331

```

```

ggplot(wisc.pr$x) +
  aes(PC1, PC2, col = diagnosis) +
  geom_point()

```

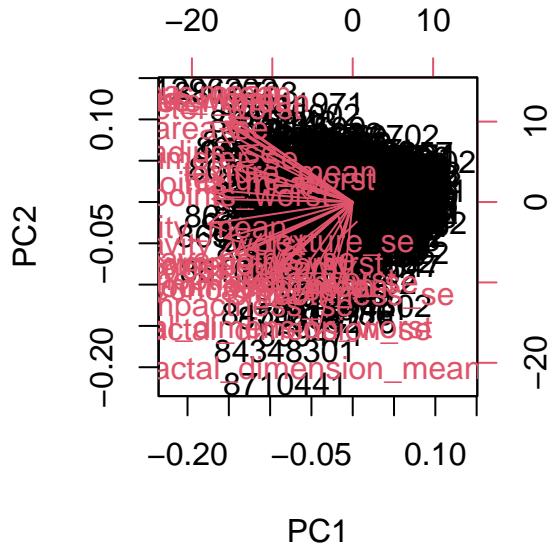


This plot clearly shows two distinct population of patients within our data set. We are going to compare this plot with some other plots.

A common visualization for PCA results is the so-called biplot.

Create a biplot of the wisc.pr using the biplot() function.

```
biplot(wisc.pr)
```



Q7. What stands out to you about this plot? Is it easy or difficult to understand? Why?

No thing really stands out about this plot. It is very messy and difficult to understand because the text labels are all clumped together

Q8 Generate a similar plot for principal components 1 and 3. What do you notice about these plots?

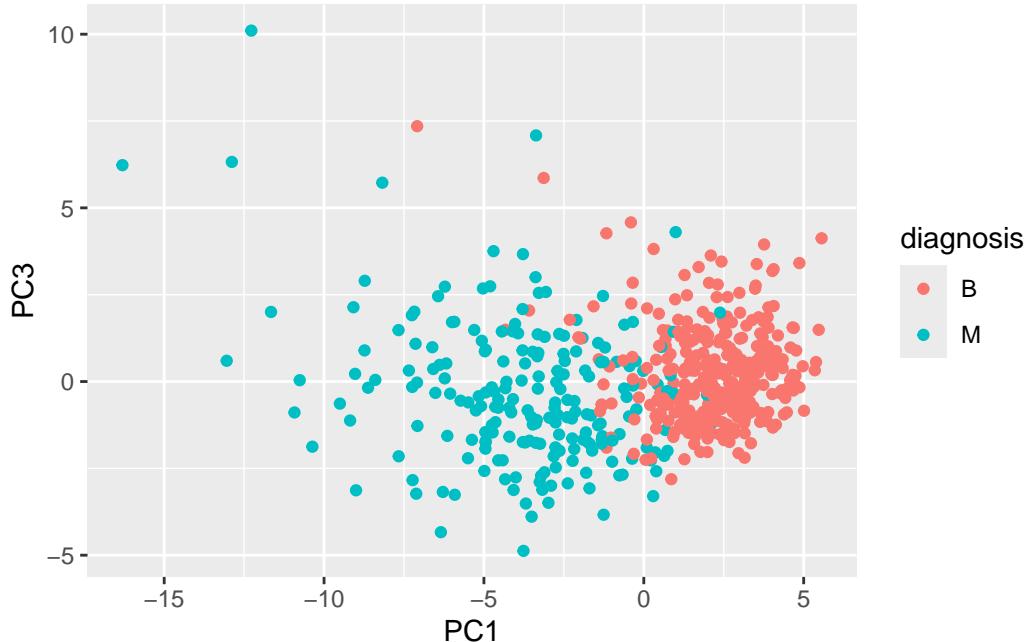
I noticed that from these scatter plot that there are clearly two distinct patient population within the dataset, and these groups can be separated by the diagnosis. Compare PC1 vs PC2 and PC1 vs PC3, we can see that PC1 vs PC2 gives a clearer separation (makes sense because PC1 and PC2 collectively captures more variance)

```
library(ggplot2)  
head(wisc.pr$x)
```

	PC1	PC2	PC3	PC4	PC5	PC6
842302	-9.184755	-1.946870	-1.1221788	3.6305364	1.1940595	1.41018364
842517	-2.385703	3.764859	-0.5288274	1.1172808	-0.6212284	0.02863116
84300903	-5.728855	1.074229	-0.5512625	0.9112808	0.1769302	0.54097615
84348301	-7.116691	-10.266556	-3.2299475	0.1524129	2.9582754	3.05073750
84358402	-3.931842	1.946359	1.3885450	2.9380542	-0.5462667	-1.22541641

843786	-2.378155	-3.946456	-2.9322967	0.9402096	1.0551135	-0.45064213
	PC7	PC8	PC9	PC10	PC11	PC12
842302	2.15747152	0.39805698	-0.15698023	-0.8766305	-0.2627243	-0.8582593
842517	0.01334635	-0.24077660	-0.71127897	1.1060218	-0.8124048	0.1577838
84300903	-0.66757908	-0.09728813	0.02404449	0.4538760	0.6050715	0.1242777
84348301	1.42865363	-1.05863376	-1.40420412	-1.1159933	1.1505012	1.0104267
84358402	-0.93538950	-0.63581661	-0.26357355	0.3773724	-0.6507870	-0.1104183
843786	0.49001396	0.16529843	-0.13335576	-0.5299649	-0.1096698	0.0813699
	PC13	PC14	PC15	PC16	PC17	
842302	0.10329677	-0.690196797	0.601264078	0.74446075	-0.26523740	
842517	-0.94269981	-0.652900844	-0.008966977	-0.64823831	-0.01719707	
84300903	-0.41026561	0.016665095	-0.482994760	0.32482472	0.19075064	
84348301	-0.93245070	-0.486988399	0.168699395	0.05132509	0.48220960	
84358402	0.38760691	-0.538706543	-0.310046684	-0.15247165	0.13302526	
843786	-0.02625135	0.003133944	-0.178447576	-0.01270566	0.19671335	
	PC18	PC19	PC20	PC21	PC22	
842302	-0.54907956	0.1336499	0.34526111	0.096430045	-0.06878939	
842517	0.31801756	-0.2473470	-0.11403274	-0.077259494	0.09449530	
84300903	-0.08789759	-0.3922812	-0.20435242	0.310793246	0.06025601	
84348301	-0.03584323	-0.0267241	-0.46432511	0.433811661	0.20308706	
84358402	-0.01869779	0.4610302	0.06543782	-0.116442469	0.01763433	
843786	-0.29727706	-0.1297265	-0.07117453	-0.002400178	0.10108043	
	PC23	PC24	PC25	PC26	PC27	
842302	0.08444429	0.175102213	0.150887294	-0.201326305	-0.25236294	
842517	-0.21752666	-0.011280193	0.170360355	-0.041092627	0.18111081	
84300903	-0.07422581	-0.102671419	-0.171007656	0.004731249	0.04952586	
84348301	-0.12399554	-0.153294780	-0.077427574	-0.274982822	0.18330078	
84358402	0.13933105	0.005327110	-0.003059371	0.039219780	0.03213957	
843786	0.03344819	-0.002837749	-0.122282765	-0.030272333	-0.08438081	
	PC28	PC29	PC30			
842302	-0.0338846387	0.045607590	0.0471277407			
842517	0.0325955021	-0.005682424	0.0018662342			
84300903	0.0469844833	0.003143131	-0.0007498749			
84348301	0.0424469831	-0.069233868	0.0199198881			
84358402	-0.0347556386	0.005033481	-0.0211951203			
843786	0.0007296587	-0.019703996	-0.0034564331			

```
ggplot(wisc.pr$x) +
  aes(PC1, PC3, col = diagnosis) +
  geom_point()
```



## Variance Explained

Calculate the variance of each principal component by squaring the sdev component of wisc.pr (i.e. `wisc.pr$sdev^2`). Save the result as an object called pr.var.

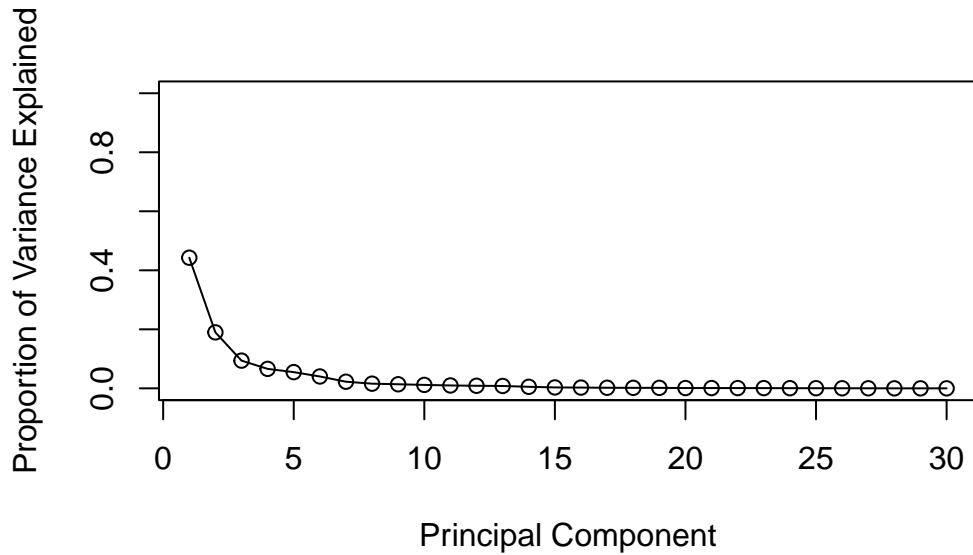
```
# Calculate variance of each component
pr.var <- wisc.pr$sdev^2
head(pr.var)
```

```
[1] 13.281608 5.691355 2.817949 1.980640 1.648731 1.207357
```

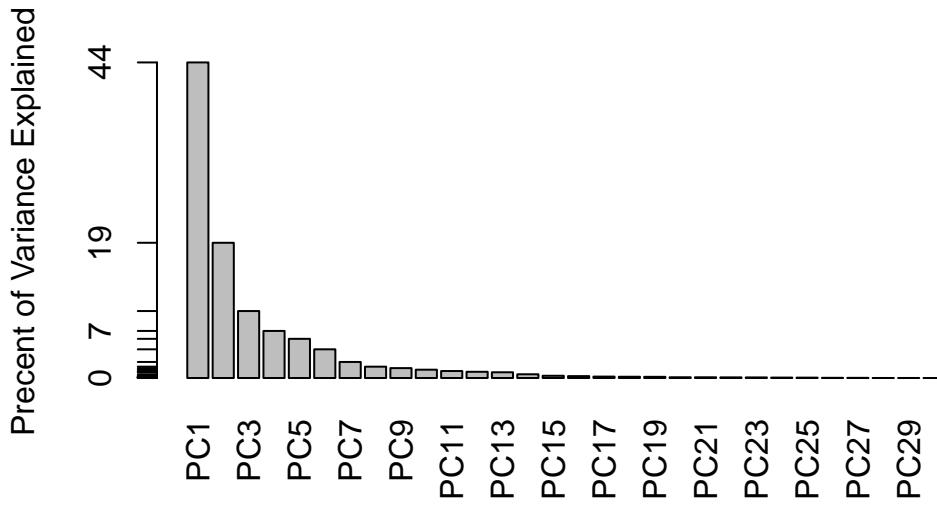
Calculate the variance explained by each principal component by dividing by the total variance explained of all principal components. Assign this to a variable called pve and create a plot of variance explained for each principal component.

```
# Variance explained by each principal component: pve
pve <- pr.var / sum(pr.var)

# Plot variance explained for each principal component
plot(pve, xlab = "Principal Component",
      ylab = "Proportion of Variance Explained",
      ylim = c(0, 1), type = "o")
```



```
# Alternative scree plot of the same data, note data driven y-axis
barplot(pve, ylab = "Percent of Variance Explained",
         names.arg=paste0("PC",1:length(pve)), las=2, axes = FALSE)
axis(2, at=pve, labels=round(pve,2)*100 )
```

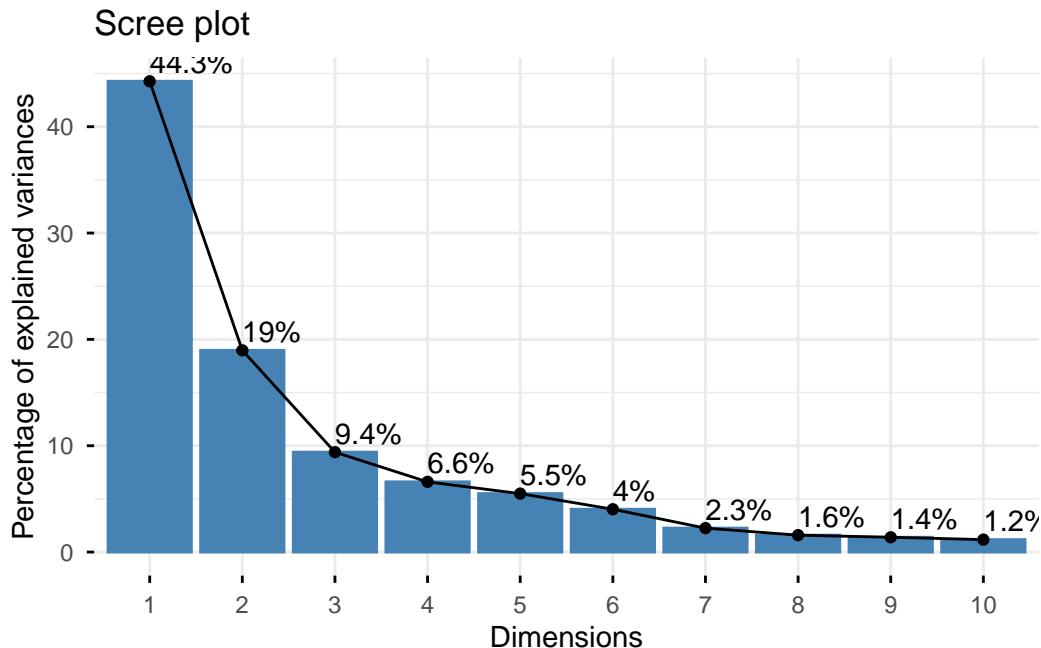


```
## ggplot based graph
#install.packages("factoextra")
library(factoextra)
```

Welcome! Want to learn more? See two factoextra-related books at <https://goo.gl/ve3WBa>

```
fviz_eig(wisc.pr, addlabels = TRUE)
```

Warning in geom\_bar(stat = "identity", fill = barfill, color = barcolor, :  
Ignoring empty aesthetic: `width`.



### Communicating PCA results

Q9. For the first principal component, what is the component of the loading vector (i.e. `wisc.pr$rotation[,1]`) for the feature `concave.points_mean`?

The loading vector of the first PC for the feature `concave.point_mean` is -0.2608538

Q10. What is the minimum number of principal components required to explain 80% of the variance of the data?

5

```
summary(wisc.pr)
```

Importance of components:

	PC1	PC2	PC3	PC4	PC5	PC6	PC7
Standard deviation	3.6444	2.3857	1.67867	1.40735	1.28403	1.09880	0.82172
Proportion of Variance	0.4427	0.1897	0.09393	0.06602	0.05496	0.04025	0.02251
Cumulative Proportion	0.4427	0.6324	0.72636	0.79239	0.84734	0.88759	0.91010
	PC8	PC9	PC10	PC11	PC12	PC13	PC14
Standard deviation	0.69037	0.6457	0.59219	0.5421	0.51104	0.49128	0.39624
Proportion of Variance	0.01589	0.0139	0.01169	0.0098	0.00871	0.00805	0.00523

```

Cumulative Proportion 0.92598 0.9399 0.95157 0.9614 0.97007 0.97812 0.98335
                  PC15    PC16    PC17    PC18    PC19    PC20    PC21
Standard deviation    0.30681 0.28260 0.24372 0.22939 0.22244 0.17652 0.1731
Proportion of Variance 0.00314 0.00266 0.00198 0.00175 0.00165 0.00104 0.0010
Cumulative Proportion 0.98649 0.98915 0.99113 0.99288 0.99453 0.99557 0.9966
                  PC22    PC23    PC24    PC25    PC26    PC27    PC28
Standard deviation    0.16565 0.15602 0.1344 0.12442 0.09043 0.08307 0.03987
Proportion of Variance 0.00091 0.00081 0.0006 0.00052 0.00027 0.00023 0.00005
Cumulative Proportion 0.99749 0.99830 0.9989 0.99942 0.99969 0.99992 0.99997
                  PC29    PC30
Standard deviation    0.02736 0.01153
Proportion of Variance 0.00002 0.00000
Cumulative Proportion 1.00000 1.00000

```

## Hierarchical clustering

Just clustering the original data is not very helpful, so...

First scale the wisc.data data and assign the result to data.scaled.

```
# Scale the wisc.data data using the "scale()" function
data.scaled <- scale(wisc.data)
```

Calculate the (Euclidean) distances between all pairs of observations in the new scaled dataset and assign the result to data.dist.

```
data.dist <- dist(data.scaled)
```

Create a hierarchical clustering model using complete linkage. Manually specify the method argument to hclust() and assign the results to wisc.hclust.

```
wisc.hclust <- hclust(data.dist,method = "complete")
```

## Results of hierarchical clustering

Q11. Using the plot() and abline() functions, what is the height at which the clustering model has 4 clusters?

At height = 19, we can get 4 clusters

- Note: But this is not very helpful because it really doesn't relay any insight/information about the dataset to us. So we give up on this approach

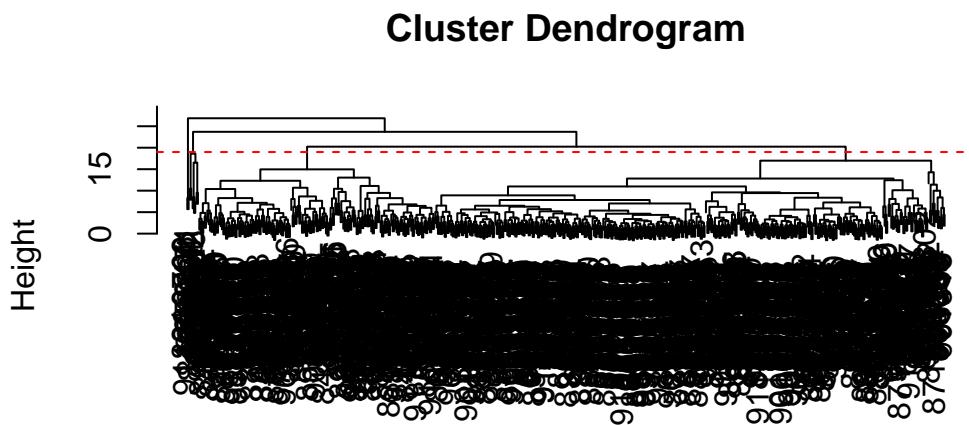
```
plot(wisc.hclust)
head(cutree(wisc.hclust, k=4))
```

```
842302 842517 84300903 84348301 84358402 843786
1       1       1       2       1       1
```

```
table(cutree(wisc.hclust,k=4))
```

1	2	3	4
177	7	383	2

```
abline(h=19, col="red", lty=2)
```



```
data.dist
hclust (*, "complete")
```

## Selecting number of clusters

Next we are going to compare the outputs from the hierarchical clustering model to the actual diagnosis.

Recall, we have the diagnosis factor, maybe it will help if we see where the malignant and benign diagnosis are in the h-clustering?

```
table((cutree(wisc.hclust,k=4)), diagnosis)
```

```
diagnosis
      B   M
1  12 165
2    2   5
3 343  40
4    0   2
```

Next, we are gonna explore how different number of clusters affect the ability of the hierarchical clustering to separate the different diagnosis

Q12. Can you find a better cluster vs diagnoses match by cutting into a different number of clusters between 2 and 10?

It is difficult to find a cluster vs diagnoses match that is meaningful. In comparison, it is slightly better when we cluster the data in 4, 5, 6, or 7 clusters as in these cases, most malignant diagnosis is clustered under cluster 1 and most benign diagnosis is clustered under cluster 3, but these clustering methods also generates a number of cluster that is meaningless (can be interpreted as another cluster's equivalent), so without prior knowledge of the diagnosis, they would not be meaningful to us.

```
test_cluster <- function(x){
  table((cutree(wisc.hclust,k=x)), diagnosis)
}
test_cluster(2)
```

```
diagnosis
      B   M
1 357 210
2    0   2
```

```
test_cluster(3)
```

```
diagnosis
    B   M
1 355 205
2   2   5
3   0   2
```

```
test_cluster(5)
```

```
diagnosis
    B   M
1 12 165
2   0   5
3 343 40
4   2   0
5   0   2
```

```
test_cluster(6)
```

```
diagnosis
    B   M
1 12 165
2   0   5
3 331 39
4   2   0
5 12   1
6   0   2
```

```
test_cluster(7)
```

```
diagnosis
    B   M
1 12 165
2   0   3
3 331 39
4   2   0
5 12   1
6   0   2
7   0   2
```

```
test_cluster(8)
```

```
diagnosis
    B   M
1 12 86
2 0 79
3 0 3
4 331 39
5 2 0
6 12 1
7 0 2
8 0 2
```

```
test_cluster(9)
```

```
diagnosis
    B   M
1 12 86
2 0 79
3 0 3
4 331 39
5 2 0
6 12 0
7 0 2
8 0 2
9 0 1
```

```
test_cluster(10)
```

```
diagnosis
    B   M
1 12 86
2 0 59
3 0 3
4 331 39
5 0 20
6 2 0
7 12 0
8 0 2
9 0 2
10 0 1
```

This confirms that this clustering is useless.

## Using different methods

Q13. Which method gives your favorite results for the same data.dist dataset?  
Explain your reasoning

the complete linkage method usually gives me my favorite results for the same data.dist dataset because it always give me distinct populations/clustering by using the largest of all pair-wise similarity to determine the clustering.

## Optional: K-means Clustering

### k-means clustering and comparing results

We will next try create a k-means clustering model on the Wisconsin breast cancer data and compare the results to the acutual diagnosis

First, we will create a k-means model on `wisc.data`

```
wisc.km <- kmeans(data.scaled, centers=2, nstart=3)
table(wisc.km$cluster, diagnosis)
```

```
diagnosis
      B    M
1 343  37
2 14   175
```

Q14. How well does k-means separate the two diagnoses?How does it compare to your hclust results?

The k-means separate the two diagnoses fairlyl well, with the malignant and benign diagnoses in cluster 1 and 2 respectively. It is does do a better job compare to the hclust we performed earlier.

```
test_cluster(4)
```

```

diagnosis
  B   M
1 12 165
2  2   5
3 343  40
4  0   2

```

## Combining methods

So let's try combine PCA with h-clustering since clustering of the origianl data is not productive, PCA is promising, so maybe we can cluster from PCA results (clustering in PC space)

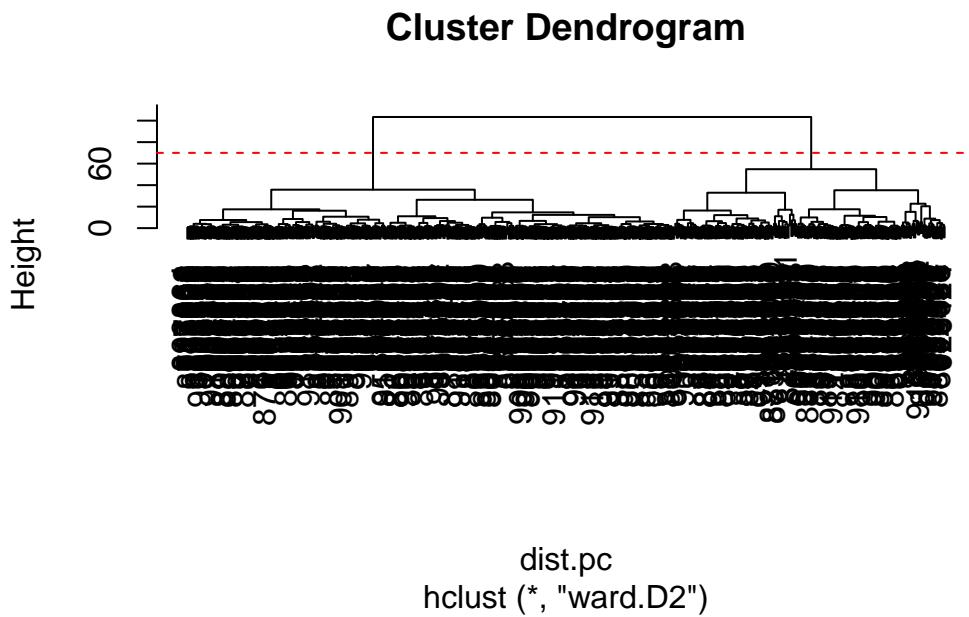
### Clustering on PCA results

Since the first 3 PCs can capture 70% of variance, we can try to cluster them first

```

# First we get the frist 3 PCs, pastt it through a dist matrix function, call it dist.pc
dist.pc <- dist(wisc.pr$x[, 1:3])
wisc.pr.hclust <- hclust(dist.pc, method="ward.D2")
plot(wisc.pr.hclust)
abline(h=70, col="red", lty=2)

```



It is good that we can see there's two main clusters. Now we are gonna try to see if this clustering is good by comparing it with the diagnosis results.

To get our clustering membership vector, we cut the tree at a desired height to yeild 2 cluster (with `cutree()` at either  $h=70$ , or  $k=2$ ) We can the use `table()` to compare the clustering results with the diagnosis results. This clustering groups is fairly accurate when compare to expert diagnosis. (For Diagnosis purpose, for malignant, we have 24 false positives out of 203 cases.)

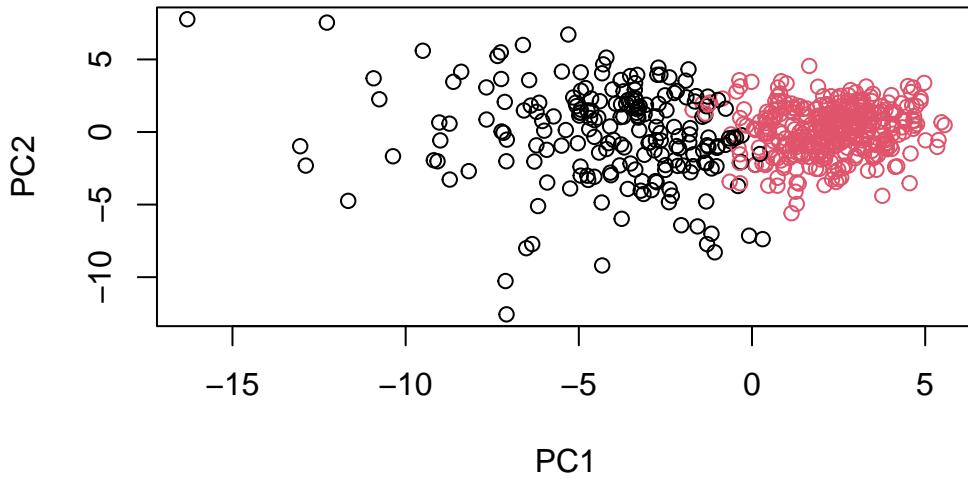
```
grps <- cutree(wisc.pr.hclust, k=2)
table(grps)
```

```
grps
 1   2
203 366
```

```
table(grps, diagnosis)
```

```
diagnosis
grps   B    M
 1  24 179
 2 333  33
```

```
plot(wisc.pr$x[,1:2], col=grps)
```



## Sensitivity & Specificity

Q15. How well does the newly created model with four clusters separates out the two diagnosis

The newly created model with 2 clusters separates out eh two diagnosis pretty neatly, as most of the malignant diagnosis is in cluster 1 while most of the benign diagnosis is in cluster 2

Q16. How well do the k-means and hierarchical clustering you created in previous section do in terms of separating the diagnoses?

The k-means clustering does a decent job separating the diagnoses (not as well as PCA); the hierarchical clustering does not do a goo job separating the diagnosis.

```
table(wisc.km$cluster,diagnosis)
```

diagnosis	
	B M
1	343 37
2	14 175

```
test_cluster(4)
```

```
diagnosis
  B   M
1 12 165
2 2   5
3 343 40
4 0   2
```

Sensitivity refers to a test's ability to correctly detect ill patients who do have the condition. In our example here the sensitivity is the total number of samples in the cluster identified as predominantly malignant (cancerous) divided by the total number of known malignant samples. In other words:  $TP/(TP+FN)$ .

Specificity relates to a test's ability to correctly reject healthy patients without a condition. In our example specificity is the proportion of benign (not cancerous) samples in the cluster identified as predominantly benign that are known to be benign. In other words:  $TN/(TN+FN)$ .

Now let's try calculate these numbers from the matrix

```
N <- table(grps, diagnosis)
N
```

```
diagnosis
grps   B   M
1 24 179
2 333 33
```

```
Sensitivity <- N[1,2]/sum(N[1,])
Sensitivity
```

```
[1] 0.8817734
```

```
Specificity <- N[2,1]/sum(N[2,])
Specificity
```

```
[1] 0.9098361
```

Q17. Which of your analysis procedure resulted in a clustering model with the best specificity? How about sensitivity?

Both K-means and PCA-hclust combined clustering both results in the highest specificity. The hierarchical clustering with k=2 has the highest sensitivity, but this is because almost all patients are grouped under cluster 1, so it has good ability to correctly detect ill patients who do have the condition. Otherwise, the sensitivities of the k-means method and the combined method are the same.

```
N <- table(grps, diagnosis)
N
```

	diagnosis	
grps	B	M
1	24	179
2	333	33

```
Sensitivity <- N[1,2]/sum(N[1,])
Sensitivity
```

```
[1] 0.8817734
```

```
Specificity <- N[2,1]/sum(N[2,])
Specificity
```

```
[1] 0.9098361
```

```
M <- table(wisc.km$cluster, diagnosis)
M
```

	diagnosis	
	B	M
1	343	37
2	14	175

```
Sensitivity <- N[1,2]/sum(N[1,])
Sensitivity
```

```
[1] 0.8817734
```

```
Specificity <- N[2,1]/sum(N[,])  
Specificity
```

```
[1] 0.9098361
```

```
0 <- table(cutree(wisc.hclust,k=2), diagnosis)  
0
```

```
diagnosis  
      B    M  
1 357 210  
2   0   2
```

```
Sensitivity <- 0[1,2]/sum(N[,])  
Sensitivity
```

```
[1] 1.034483
```

```
Specificity <- 0[2,1]/sum(N[,])  
Specificity
```

```
[1] 0
```

## Prediction

We can use our PCA model for predicting with new input patient samples

```
#url <- "new_samples.csv"  
url <- "https://tinyurl.com/new-samples-CSV"  
new <- read.csv(url)  
npc <- predict(wisc.pr, newdata=new)  
npc
```

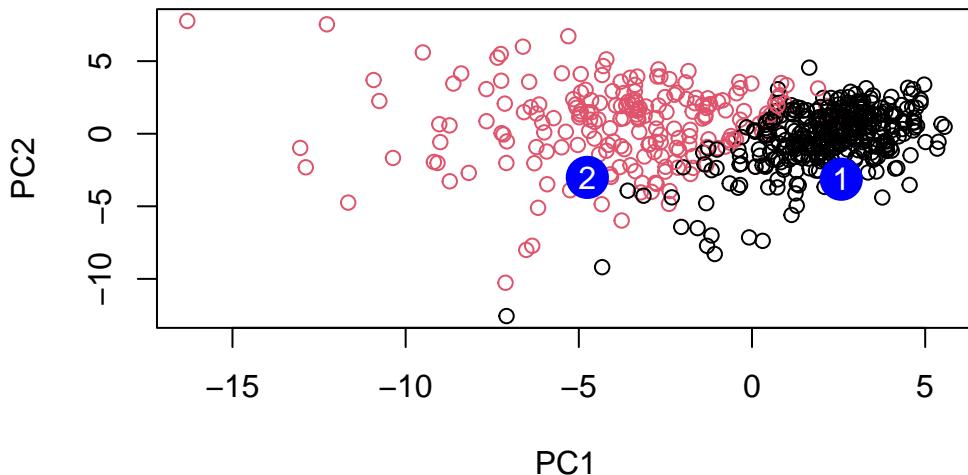
	PC1	PC2	PC3	PC4	PC5	PC6	PC7
[1,]	2.576616	-3.135913	1.3990492	-0.7631950	2.781648	-0.8150185	-0.3959098
[2,]	-4.754928	-3.009033	-0.1660946	-0.6052952	-1.140698	-1.2189945	0.8193031
	PC8	PC9	PC10	PC11	PC12	PC13	PC14

```
[1,] -0.2307350 0.1029569 -0.9272861 0.3411457 0.375921 0.1610764 1.187882
[2,] -0.3307423 0.5281896 -0.4855301 0.7173233 -1.185917 0.5893856 0.303029
PC15 PC16 PC17 PC18 PC19 PC20
[1,] 0.3216974 -0.1743616 -0.07875393 -0.11207028 -0.08802955 -0.2495216
[2,] 0.1299153 0.1448061 -0.40509706 0.06565549 0.25591230 -0.4289500
PC21 PC22 PC23 PC24 PC25 PC26
[1,] 0.1228233 0.09358453 0.08347651 0.1223396 0.02124121 0.078884581
[2,] -0.1224776 0.01732146 0.06316631 -0.2338618 -0.20755948 -0.009833238
PC27 PC28 PC29 PC30
[1,] 0.220199544 -0.02946023 -0.015620933 0.005269029
[2,] -0.001134152 0.09638361 0.002795349 -0.019015820
```

```
plot(wisc.pr$x[,1:2], col=diagnosis)
cat(paste(levels(diagnosis), "=", palette() [seq_along(levels(diagnosis))]), collapse = ", "))

B = black,M = #DF536B
```

```
points(npc[,1], npc[,2], col="blue", pch=16, cex=3)
text(npc[,1], npc[,2], c(1,2), col="white")
```



Q18. Which of these new patients should we prioritize for follow up based on your results?

We should prioritize patient 2 based on the results since patient 2 is placed in the cluster with malignant diagnosis.